

# Exporting Splunk Data at Scale with Scribl

## scribl.py

Exporting Splunk Data at Scale with [Scribl](#). This is a python script that can be run on each Splunk Indexer for the purpose of exporting historical bucket data (raw events + metadata) at scale by balancing the work across multiple CPUs then forwarding to Cribl.

## Demo Videos

- [HD - Exporting Splunk Data at Scale with Scribl](#)
- [4K - Exporting Splunk Data at Scale with Scribl](#)

## Background

Splunk to Cribl = [scribl](#) (#thanksKam)

Exporting large amounts of previously indexed data from Splunk is challenging via the Splunk-supported approaches detailed here: <https://docs.splunk.com/Documentation/Splunk/8.2.6/Search/Exportsearchresults>.


The core Splunk binary in every install provides a switch (cmd exporttool) that allows you to export the data from the compressed buckets on the indexers back into their original raw events. You can dump them to very large local csv files or stream them to stdout so a script can redirect over the network to a receiver such as Cribl Stream. This switch has been used by others for quite a while but it isn't well documented.

Assuming that Splunk is installed in /opt/splunk/, the below commands can be applied to a particular bucket in an index called "bots" to export it.

Exporting to stdout:

```
 /opt/splunk/bin/splunk cmd exporttool /opt/splunk/var/lib/splunk/bots/db/db_1564739504_1564732800_2394 /dev/stdout -csv
```

Exporting to a local csv file:

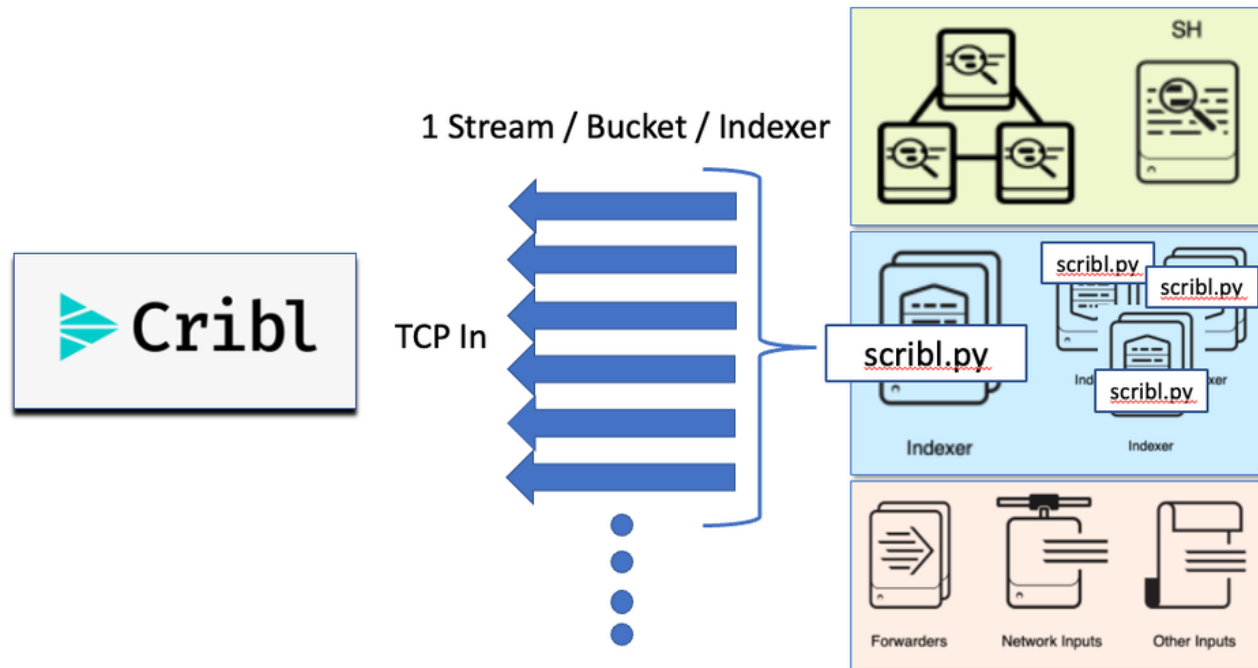
```
 /opt/splunk/bin/splunk cmd exporttool /opt/splunk/var/lib/splunk/bots/db/db_1564739504_1564732800_2394  
/exports/bots/db_1564739504_1564732800_2394.csv -csv
```

There will be many buckets so some poor soul will need to build a script to export all or some of the buckets and some sort of parallelization should be used to speed the process up. The exported data will be very large (uncompressed, 3-20x) compared to the size of the individual buckets that make up the index!

## Requirements

Splunk stores its collected data on the indexers within the "Indexing Tier" as detailed below. The data is compressed and stored in a collection of time series buckets that reside on each indexer. Each bucket contains a rawdata journal, along with associated tsidx, and metadata files. The search heads access these buckets and it's very rare for someone to access them directly from the indexer CLI unless

there is a need to export data to retrieve the original raw events. We will use the indexer CLI to export the original raw events (per bucket and in parallel) as well as a few other pieces of important metadata as detailed below.



For a deeper dive into how Splunk indexes data, see this: [🔴 How the indexer stores indexes - Splunk Documentation](#)

You will need:

- CLI access to each Linux indexer with the index/buckets that need to be exported. This process only applies to on-prem or non-SplunkCloud deployments.
- To install nc (netcat) on each indexer to act as the transport mechanism until we have enough demand to build the transport into the script.
- To make sure outbound communication from each indexer to the Cribl Worker TCP port is open.

## Frozen Data

The Splunk exporttool switch that scribl depends on requires a complete hot/warm/cold directory containing all of the metadata files in addition to the journal.gz file. When buckets are moved to a frozen archive, all of the metadata files are removed with only the journal.gz file remaining. Scribl can not extract raw events from frozen archives.

Buckets must first be “thawed” as described [here](#). It’s a straightforward process of copying the frozen buckets somewhere and running a “splunk rebuild” for each bucket to recreate the metadata. Scribl can be run against this thawed data.

## Technical

### Scale

We achieve scale for large volumes of data by processing buckets in parallel across as many CPUs as you would like to dedicate AND by streaming the data directly from disk with a single read to Cribl without ever having to write extracted/uncompressed event data to disk. Extracting/uncompressing the event data to disk would result in enormous disk IO bottlenecks and disk space consumption.

Disk speed (IOPS) and the number of CPUs are generally your limiting factors on the indexers. While disk speed is a factor, it's usually not a factor in the overall scale picture because Splunk indexers will usually have high IOPS capabilities. You can monitor the linux processes to get a feel for whether scribl processes are in SLEEP or RUN mode. If they spend the majority of their time in SLEEP mode, they are being throttled by disk, network, Cribl workers, etc and adding more CPUs will probably not buy you more speed.

The scribl script running on your indexers and Cribl Stream workers are built to scale and will usually not be your bottleneck. Your bottlenecks will almost certainly be bandwidth constraints between your indexers and your final destination. Depending on where you deploy your Cribl Stream workers, that bandwidth bottleneck might exist between the indexers and Cribl workers or between your Cribl workers and the final destination. If you happen to have unlimited bandwidth, you will likely find your next bottleneck to be the ingest rate at your destination platform.

## Exported Data Format

The exported data will be csv formatted with a header followed by the individual events. It's important to call out that these events are often multiline events with the most common example being windows logs. The below events are examples that are generated by Splunk and then passed via stdin to the scribl.py script.

The `_raw` field contains the original event and the other fields were captured/created during ingest. `_time` is the time extracted from the event which will be the primary time reference used by the destination analytics platform. The `sourcetype` field will likely be what is used by the destination to determine how to parse and where to route the event.

Example:

• `"_time",source,host,sourcetype,"_raw","_meta"`

```
1564734905,"source::10.1.1.1","host::hogshead","sourcetype::fgt_utm","date=2019-08-02 time=08:35:05 devname=hogshead
devid=FGT60D4614044725 logid=1059028704 type=utm subtype=app-ctrl eventtype=app-ctrl-all level=information vd=root
appid=38131 user="" srcip=10.1.1.103 srcport=51971 srcintf="" internal"" dstip=172.217.11.227 dstport=443 dstintf="" wan1""
profilename="" applist="" proto=6 service="" HTTPS"" policyid=1 sessionid=594789 applist="" default"" appcat="" General.Interest""
app="" Google.Accounts"" action=pass hostname="" ssl.gstatic.com"" url="" msg="" General.Interest: Google.Accounts,""
apprisk=elevated","_indextime::1564734907 _subsecond::000 syslog-server::jupiter severity::notice facility::user punct::""=--
_:: = = = = = - - - = = = = = \"" = ... = \"" = ... = ""
```

```
1564734846,"source::WinEventLog:Microsoft-Windows-
PowerShell/Operational","host::titan","sourcetype::XmlWinEventLog:Microsoft-Windows-PowerShell/Operational","Event
xmlns='http://schemas.microsoft.com/win/2004/08/events/event'><System><Provider Name='Microsoft-Windows-PowerShell'
Guid='{A0C1853B-5C40-4B15-8766-3CF1C58F985A}'><EventID>4103</EventID><Version>1</Version><Level>4</Level>
<Task>106</Task><Opcode>20</Opcode><Keywords>0x0</Keywords><TimeCreated SystemTime='2019-08-
02T08:34:06.167139700Z'/><EventRecordID>5968761</EventRecordID><Correlation ActivityID='{135BC459-4718-0000-AAD1-
74131847D501}'><Execution ProcessID='3016' ThreadID='3720'/><Channel>Microsoft-Windows-
PowerShell/Operational</Channel><Computer>titan.thirstyberner.com</Computer><Security UserID='S-1-5-18'/></System>
<EventData><Data Name='ContextInfo'>          Severity = Informational
```

Host Name = Default Host

Host Version = 4.0

Host ID = 89297903-4c6b-4e9d-b0a4-49c76b2c36ae


Engine Version = 4.0

Runspace ID = 657f43dd-6fb5-42c9-8b93-154f3a1e53dd

Pipeline ID = 1

Command Name = Out-Null

```
</Data></EventData></Event>"", "_indextime::1564734908 punct::< _='://../////"><>< _='--' _='{----}'><></><></><><"
```

 `./scribl.py -h`

usage: `scribl.py [-h] [-t] [-n NUMSTREAMS] [-l LOGFILE] [-et EARLIEST] [-lt LATEST] [-d DIRECTORY] [-r REMOTEIP] [-p REMOTEPORT]`

This is to be run on a Splunk Indexer for the purpose of exporting buckets and streaming their contents to Cribl Stream

optional arguments:

- `-h, --help` show this help message and exit
- `-t, --TLS` Send with TLS enabled
- `-n NUMSTREAMS, --numstreams NUMSTREAMS`  
the number of parallel stream to utilize
- `-l LOGFILE, --logfile LOGFILE`  
specify the location to write/append the logging
- `-et EARLIEST, --earliest EARLIEST`  
specify the earliest epoch time for bucket selection
- `-lt LATEST, --latest LATEST`  
specify the latest epoch time for bucket selection

required named arguments:

- `-d DIRECTORY, --directory DIRECTORY`  
Source directory containing the buckets

```
-r REMOTEIP, --remoteIP REMOTEIP
    Remote address to send the exported data to
-p REMOTEPORT, --remotePort REMOTEPORT
    Remote TCP port to be used
```

```
scribl.py -d /opt/splunk/var/lib/splunk/bots/db/ -r 14.2.39.121 -p 20000 -t n4 -l /tmp/scribl.log -et 1564819155 -lt 1566429310
```

## On No! My Splunk license expired!

Worry not, my friend. When the enterprise license expires, Splunk customers are free to use the 60-day trial or even the free version of Splunk to perform the export. Sanity check my claim here:

<https://docs.splunk.com/Documentation/Splunk/9.0.0/Admin/MoreaboutSplunkFree> .

We don't care about indexing new data and we don't care about distributed search since we will use the trial/free Splunk binary in a standalone manner on each of the indexers that have data we need to migrate. Just install trial/free Splunk on top of or alongside the existing install and point scribl.py at your splunk binary and the directory containing the buckets you need to export.

## Cribl Stream Config

You can get started instantly with Cribl Cloud or even using the Cribl Free [license option](#) but keep in mind daily ingest limits (very generous) and # of cores (also very generous at 10) that can be used may factor into a full scale data export. If you choose to install Cribl Stream on-prem in your own cloud, the [documentation](#) is your friend and will get you going quickly.

Once you have satisfied the above requirements (CLI, nc, and firewall) on your Splunk indexers, grab the [scribl.py script from the github repo](#) and copy it over to each indexer. The only thing in the script that is hard coded is the default install location of Splunk (/opt/splunk) which you can easily modify if you are running a non-default config. Keep in mind that we are running the script directly on the Splunk indexers and a python binary is kept under \$SPLUNK\_HOME/bin.

Add a new Source:

The screenshot shows the Cribl Stream configuration page for a source named 'scribl'. The breadcrumb trail at the top is 'Stream > Sources > TCP > scribl'. Below this is a navigation bar with tabs: 'Configure' (selected), 'Status', 'Charts', 'Live Data', and 'Logs'. A 'Help' link with a question mark icon is on the right. On the left side, there is a sidebar menu with the following items: 'General Settings' (highlighted), 'TLS Settings (Server Side)', 'Persistent Queue Settings', 'Processing Settings' (with an expand/collapse arrow), 'Custom Command', 'Event Breakers', 'Fields', 'Pre-Processing', 'Advanced Settings', and 'Connected Destinations'. The main configuration area on the right contains the following fields and controls:

- Input ID\***: A text field containing 'scribl'. To its right is an 'Enabled' toggle switch set to 'Yes'.
- \_\_inputId**: A text field containing 'tcp:scribl' with a copy icon to its right.
- Address\***: A text field containing '0.0.0.0'.
- Port\***: A text field containing '20000'.
- Enable Header**: A toggle switch set to 'No'.
- Tags**: A text field containing 'scribl' with a close icon to its right.

Enable and Config TLS (recommended):

Stream > Groups > default > Sources > TCP > scribl

Configure   Status   Charts   Live Data   Logs

General Settings

TLS Settings (Server Side)

Processing Settings ^

Event Breakers

Fields

Pre-Processing

Advanced Settings

Connected Destinations

Enabled ⓘ  

Yes ☒ Autofill?

Certificate name ⓘ  

Select one

Private key path\* ⓘ  

/opt/criblcerts/criblcloud.key

Passphrase ⓘ  

Enter passphrase

Certificate path\* ⓘ  

/opt/criblcerts/criblcloud.crt

CA certificate path ⓘ  

Enter CA certificate path

Authenticate client (mutual auth) ⓘ ☐ No

Minimum TLS version ⓘ  

TLSv1.2

Maximum TLS version ⓘ  

Select one

Create a new Event Breaker and set it to "[\n\r]+\d{10}," as detailed below:

Rule Name\* ?

scribl

In Out

[Upload Sample File](#)

Filter Condition\* ?

true



## EVENT BREAKER SETTINGS

Enabled ? ☒ Yes

Event Breaker Type\* ?

Regex

Event Breaker\* ?

/ [\n\r]+\d{10},



Max Event Bytes ?

51200

## TIMESTAMP SETTINGS

Timestamp Anchor\* ?



Timestamp Format\* ?

☒ Autotimestamp Scan Depth ? 150☐ Manual Format ?☐ Current Time ?

Default Timezone ?

Local

Earliest timestamp allowed ?

-420weeks

Future timestamp allowed ?

+1week

## ADD FIELDS TO EVENTS ?

+ Add Field

Paste your events here or upload a sample file

☒ Timestamp Anchor ☐ Event Breaker ☐ Timestamp

Cancel

OK

Validate:

Rule Name\* ?

exporttool

In Out

[Upload Sample File](#)

Filter Condition\* ?

true



## EVENT BREAKER SETTINGS

Enabled ? ☒ Yes

Event Breaker Type\* ?

Regex

Event Breaker\* ?

/ [\n\r]+\d{10},



Max Event Bytes ?

1200

```
1      @ _raw:
2022-05-25      "time",source,host,sourcetype,"_raw","_meta"
15:42:08.572   # _time: 1653511328.572
-05:00         @ cribl_breaker: exporttool

2      @ _raw: "source:10.1.1.1","host:hogshead","sourcetype:fgt_utm","date=2019-08-02 time=08:35:05 devname=hogshead
2022-05-25      devid=FGT604614044725 logid=1059028704 type=utm subtype=app-ctrl eventtype=app-ctrl-all le... Show more
15:42:08.574   # _time: 1653511328.574
-05:00         @ cribl_breaker: exporttool

3      @ _raw: "source:WinEventLog:Microsoft-Windows-PowerShell/Operational","host:titan","sourcetype:XmlWinEventLog:
2022-05-25      Microsoft-Windows-PowerShell/Operational","cEvent xmlns='http://schemas.microsoft.com/win/20... Show more
15:42:08.574   # _time: 1653511328.574
-05:00         @ cribl_breaker: exporttool
```

Attach the new Event Breaker to your Source:

The screenshot shows the Splunk configuration interface for a source named 'scribl'. The left sidebar contains a menu with options: General Settings, TLS Settings (Server Side), Persistent Queue Settings, Processing Settings, Custom Command, and Event Breakers (which is highlighted). The main content area is titled 'Event Breaker rulesets' and shows a list of rulesets. A single ruleset named 'scribl' is listed with the description 'Ingest events from Splunk via the scribl script from Cribl (1 rule)'. Below the list, the 'System Default Rule' is expanded, showing filter conditions and event breaker settings. The 'Event Breaker buffer timeout' is set to 10000.

ream > Sources > TCP > scribl

onfigure Status Charts Live Data Logs Help

General Settings

TLS Settings (Server Side)

Persistent Queue Settings

Processing Settings

Custom Command

Event Breakers

Event Breaker rulesets

1 scribl Ingest events from Splunk via the scribl script from Cribl (1 rule)

System Default Rule Filter Condition: true Event Breaker: /[\n\r]+(?:\s)/ Timestamp Anchor: / ^ / Timestamp Format: Auto:150 Default Timezone: Local Max Event Bytes: 51200

+ Add ruleset

Event Breaker buffer timeout

10000

Add a new route and attach a pipeline:

The screenshot shows the Splunk configuration interface for a route named 'scribl'. The top section shows a table with columns: #, Route, Filter, Pipeline/Outp... Events, and a status column. The table has one entry with # 1, Route scribl, Filter \_\_inputId=..., Pipeline/Outp... Events scribl splunk\_hec:s..., and a status of 100.000%. Below the table, the route configuration is shown. The Route Name is 'scribl'. The Filter is '\_\_inputId='tcp:scribl''. The Pipeline is 'scribl'. The Enable Expression is 'No'. The Output is 'splunk\_hec:scribl-HEC'. The Description is 'Data arriving from the scribl.py script'. The Final checkbox is checked.

Route Filter Pipeline/Outp... Events In Out Dropped All

1 scribl \_\_inputId=... scribl splunk\_hec:s... 100.000% On

Route Name\* scribl

Filter ? \_\_inputId='tcp:scribl'

Pipeline\* ? scribl

Enable Expression ? No

Output ? splunk\_hec:scribl-HEC

Description ? Data arriving from the scribl.py script

Final ? Yes

Attach functions to your pipeline to transform the inbound data as detailed below:



1

Parser

true

On

...

Filter ?

true

Help ?

Description ?

Extract the initial fields from \_raw

Final ?

No

Operation Mode\* ?

Extract

Type\* ?

CSV

Library ?

Select from Library

Source Field ?

\_raw

Destination Field ?

Destination field name

List of Fields ?

source x host x sourcetype x \_raw x meta x

2

Regex Extract

true

On

..

Filter ?

true

Help ?

Description ?

Capture the original index time and use that as \_time

Final ?

No

Regex\* ?

/^\_indextime::(<\_time>\d{10})

/

R

Additional Regex

+ Add Regex

Source Field ?

meta

ADVANCED SETTINGS

Max Exec ?

100

Field Name Format Expression ?

Enter field name format expression.

Overwrite Existing Fields ?

Yes

2MasktrueOn

Filter ? Help ?

true

Description ?

Enter a description

Final ? No

Masking Rules\*

Match Regex ?	Replace Expression ?	
/ (?w*::)		
+ Add Rule		

Apply to Fields ?

host x source x sourcetype x

> ADVANCED SETTINGS

3Regex ExtracttrueOff

Filter ? Help ?

true

Description ?

Extract \_indextime, \_subsecond, and punt fields from meta

Final ? No

Regex\* ?

/ ^\_indextime::(?<\_indextime>\d{10}) (\s\_subsecond::(?<\_subsecond>\S{4}) /

Additional Regex

+ Add Regex

Source Field ?

meta

4EvaltrueOn

Filter ? Help ?

true

Description ?

Remove the meta field

Final ? No

Evaluate Fields ?

+ Add Field

Keep Fields ?

Enter field names

Remove Fields ?

meta x

5

Eval

true

On

...

Filter ?

true

Help ?

Description ?

Remove the cribl\_breaker field

Final ?

No

Evaluate Fields ?

+ Add Field

Keep Fields ?

Enter field names

Remove Fields ?

cribl\_breaker x

Validate the transformation then attach your destination(s) to the route:

```

1      x
2022-06-15
09:44:25.632
-05:00

raw: <Event xmlns='http://schemas.microsoft.com/win/2004/08/events/event'><System><Provider Name='Microsoft-Windows-PowerShell' Guid='{A0C1853B-5C40-4B15-8766-3CF1C58F985A}' /><EventID>4103</EventID><Version>1</Version><Level>4</Level><Task>106</Task><Opcode>20</Opcode><Keywords>0x0</Keywords><TimeCreated SystemTime='2019-08-02T09:38:04.096934100Z' /><EventRecordID>6165455</EventRecordID><Correlation ActivityID='{135BC459-4718-0000-F407-7B131847D501}' /><Execution ProcessID='5080' ThreadID='932' /><Channel>Microsoft-Windows-PowerShell/Operational</Channel><Computer>titan.thirstyberner.com</Computer><Security UserID='S-1-5-18' /></System><EventData><Data Name='ContextInfo' Severity = Informational
    Host Name = Default Host
    Host Version = 4.0
    Host ID = bc05089d-06d9-49cd-8c05-5672caf6aace
    Engine Version = 4.0
    Runspace ID = 5443b262-5049-49b7-9db1-f7995525117e
    Pipeline ID = 1
    Command Name = Out-Null
    Command Type = Cmdlet
    Script Name =
    Command Path =
    Sequence Number = 258134
    User = THIRSTYBERNER\SYSTEM
    Shell ID = Microsoft.PowerShell
  </Data><Data Name='UserData'></Data><Data Name='Payload'>ParameterBinding(Out-Null): name="InputObject"; value="True"
</Data></EventData></Event> Show less

#_time: 1655304265.632
cribl_breaker: scribl:scribl
cribl_pipe: scribl
host: titan
source: WinEventLog:Microsoft-Windows-PowerShell/Operational
sourcetype: XmlWinEventLog:Microsoft-Windows-PowerShell/Operational

```

## Caveats:

### Splunk Event Sizes

You need to pay attention to event sizes in Splunk as it pertains to the Event breaking in Cribl. As noted above in the Event Breaker screenshot, the max event size has a default setting of 51200 bytes. If you use scribl to send events into Cribl Stream larger than that,

things break. Either increase your event breaking max event size, use the Cribl Stream Pipeline to drop the large events (example: by sourcetype), or do not use scribl to export the buckets containing the large events.

Here is a quick Splunk search highlighting the large events that need to be dealt with:

New Search

Save As>Create Table ViewClose

index=bots|eval l=len(\_raw)|where 25000>l|stats count values(sourcetype) by l|sort - l

All time

✓ 31,304,827 events (before 6/23/22 4:06:41.000 PM) No Event Sampling

Job

Smart Mode

Events

Patterns

Statistics (10,000)

Visualization

20 Per Page

Format

Preview

< Prev

1

2

3

4

5

6

7

8

...

Next >

		count	values(sourcetype)
2400680	31	stream:http	
174623	1	stoq	
159433	31	stoq	
156462	31	stoq	
156297	31	stoq	
153023	31	stoq	

## Bottlenecks

As mentioned above, the bottleneck you will most likely run into will be bndwidth in your data path or ingest rate at the final destination. Anything you can do to parallelize that final write will pay dividends. For example, you may want to use Cribl Stream's Output Router to write to multiple S3 buckets based on the original Splunk Index or Sourcetype if bandwidth is not your bottleneck.

## To Do:

- ~~Add min and max times to determine which buckets should be exported (complete 6/2022)~~
- Applies this to a smartstore config (auth, encryption, data format, etc)
- Sync with Spico and other partners who would be interested
- How about this as a collector?